

black

August 6, 2024

## 0.0.1 Black Friday Dataset EDA And Feature Enaineering

## 0.0.2 Cleaning and preparing the data for model Training

# 1 Problem Statement

A retail company “ABC Private Limited” wants to understand the customer purchase behaviour (specifically, purchase amount) against various products of different categories. They have shared purchase summary of various customers for selected high volume products from last month. The data set also contains customer demographics (age, gender, marital status, city\_type, stay\_in\_current\_city), product details (product\_id and product category) and Total purchase\_amount from last month.

Now, they want to build a model to predict the purchase amount of customer against various products which will help them to create personalized offer for customers against different products.

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df_train = pd.read_csv('/content/train.csv')
```

```
[ ]: df_train
```

```
[ ]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	
...	...	...	...	...	...	...	
550063	1006033	P00372445	M	51-55	13	B	
550064	1006035	P00375436	F	26-35	1	C	
550065	1006036	P00375436	F	26-35	15	B	
550066	1006038	P00375436	F	55+	1	C	
550067	1006039	P00371644	F	46-50	0	B	

```
Stay_In_Current_City_Years  Marital_Status  Product_Category_1  \
```

0	2	0	3
1	2	0	1
2	2	0	12
3	2	0	12
4	4+	0	8
...	...	...	...
550063	1	1	20
550064	3	0	20
550065	4+	1	20
550066	2	0	20
550067	4+	1	20

	Product_Category_2	Product_Category_3	Purchase
0	NaN	NaN	8370
1	6.0	14.0	15200
2	NaN	NaN	1422
3	14.0	NaN	1057
4	NaN	NaN	7969
...	...	...	...
550063	NaN	NaN	368
550064	NaN	NaN	371
550065	NaN	NaN	137
550066	NaN	NaN	365
550067	NaN	NaN	490

[550068 rows x 12 columns]

```
[ ]: df_test = pd.read_csv('/content/test.csv')
```

```
[ ]: ## append is not working
      ##df = df_train.append(df_test)
```

```
[ ]: #The error message "'DataFrame' object has no attribute 'append'" occurs
      # because the append() method was deprecated in pandas 2.0.
      # Instead, you should use the concat() function to append data frames.
      df = pd.concat([df_train, df_test], ignore_index=True)
```

```
[ ]: df.head()
```

```
[ ]:   User_ID Product_ID Gender  Age  Occupation City_Category \
0  1000001  P00069042      F  0-17         10           A
1  1000001  P00248942      F  0-17         10           A
2  1000001  P00087842      F  0-17         10           A
3  1000001  P00085442      F  0-17         10           A
4  1000002  P00285442      M  55+         16           C
```

```
      Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
```

0	2	0	3
1	2	0	1
2	2	0	12
3	2	0	12
4	4+	0	8

	Product_Category_2	Product_Category_3	Purchase
0	NaN	NaN	8370.0
1	6.0	14.0	15200.0
2	NaN	NaN	1422.0
3	14.0	NaN	1057.0
4	NaN	NaN	7969.0

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 783667 entries, 0 to 783666
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               783667 non-null  int64
1   Product_ID                            783667 non-null  object
2   Gender                                783667 non-null  object
3   Age                                    783667 non-null  object
4   Occupation                             783667 non-null  int64
5   City_Category                         783667 non-null  object
6   Stay_In_Current_City_Years            783667 non-null  object
7   Marital_Status                        783667 non-null  int64
8   Product_Category_1                    783667 non-null  int64
9   Product_Category_2                    537685 non-null  float64
10  Product_Category_3                    237858 non-null  float64
11  Purchase                              550068 non-null  float64
dtypes: float64(3), int64(4), object(5)
memory usage: 71.7+ MB
```

```
[ ]: df.describe()
```

```
[ ]:
count      User_ID      Occupation  Marital_Status  Product_Category_1  \
mean      1.003029e+06      8.079300      0.409777      5.366196
std        1.727267e+03      6.522206      0.491793      3.878160
min         1.000001e+06      0.000000      0.000000      1.000000
25%         1.001519e+06      2.000000      0.000000      1.000000
50%         1.003075e+06      7.000000      0.000000      5.000000
75%         1.004478e+06     14.000000      1.000000      8.000000
max         1.006040e+06     20.000000      1.000000     20.000000
```

	Product_Category_2	Product_Category_3	Purchase
count	537685.000000	237858.000000	550068.000000
mean	9.844506	12.668605	9263.968713
std	5.089093	4.125510	5023.065394
min	2.000000	3.000000	12.000000
25%	5.000000	9.000000	5823.000000
50%	9.000000	14.000000	8047.000000
75%	15.000000	16.000000	12054.000000
max	18.000000	18.000000	23961.000000

## 2 Drop the unwanted columns

```
[ ]: df.drop(['User_ID'], axis =1, inplace=True)
```

```
[ ]: df.head()
```

```
[ ]:
  Product_ID Gender  Age  Occupation City_Category \
0  P00069042     F  0-17          10           A
1  P00248942     F  0-17          10           A
2  P00087842     F  0-17          10           A
3  P00085442     F  0-17          10           A
4  P00285442     M  55+          16           C
```

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0	3	
1	2	0	1	
2	2	0	12	
3	2	0	12	
4	4+	0	8	

	Product_Category_2	Product_Category_3	Purchase
0	NaN	NaN	8370.0
1	6.0	14.0	15200.0
2	NaN	NaN	1422.0
3	14.0	NaN	1057.0
4	NaN	NaN	7969.0

### 2.0.1 Convert the categorical into numerical

```
[ ]: ## Converting the gender to the numerical variable
# first is to create the dummy dataset for the gender and then assigning it to
↳ the dataset
#df['Gender']=pd.get_dummies(df['Gender'],drop_first=1)
# The next is to directly map to the gender column
df['Gender']=df['Gender'].map({'F':0,'M':1})
df.head()
```

```
[ ]: Product_ID  Gender  Age  Occupation  City_Category  \
0  P00069042      0  0-17          10             A
1  P00248942      0  0-17          10             A
2  P00087842      0  0-17          10             A
3  P00085442      0  0-17          10             A
4  P00285442      1  55+          16             C

Stay_In_Current_City_Years  Marital_Status  Product_Category_1  \
0                2                0                3
1                2                0                1
2                2                0               12
3                2                0               12
4                4+                0                8

Product_Category_2  Product_Category_3  Purchase
0                NaN                NaN    8370.0
1                6.0               14.0   15200.0
2                NaN                NaN    1422.0
3               14.0                NaN    1057.0
4                NaN                NaN    7969.0
```

```
[ ]: df['Age'].unique()
```

```
[ ]: array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
      dtype=object)
```

For Age column

```
[ ]: #pd.get_dummies(df['Age'],drop_first=True)
df['Age']=df['Age'].map({'0-17':1,'18-25':2,'26-35':3,'36-45':4,'46-50':
↳5,'51-55':6,'55+':7})
```

Label encoding can also be performed for this task *##second technique from sklearn import preprocessing*

## 2.0.2 label\_encoder object knows how to understand word labels.

```
label_encoder = preprocessing.LabelEncoder()
```

### 2.0.3 Encode labels in column 'Age'.

```
df['Age']= label_encoder.fit_transform(df['Age'])
```

```
df['Age'].unique()
```

```
[ ]: df.head()
```

```
[ ]:  Product_ID  Gender  Age  Occupation  City_Category  \
0  P00069042      0    1         10             A
1  P00248942      0    1         10             A
2  P00087842      0    1         10             A
3  P00085442      0    1         10             A
4  P00285442      1    7         16             C

    Stay_In_Current_City_Years  Marital_Status  Product_Category_1  \
0                             2                0                  3
1                             2                0                  1
2                             2                0                 12
3                             2                0                 12
4                             4+                0                  8

    Product_Category_2  Product_Category_3  Purchase
0                  NaN                  NaN    8370.0
1                  6.0                  14.0   15200.0
2                  NaN                  NaN    1422.0
3                 14.0                  NaN    1057.0
4                  NaN                  NaN    7969.0
```

```
[ ]: ##fixing categorical City_category
## As like earlier we discussed we are creating the dummies
# and adding it to the dataset
df_city=pd.get_dummies(df['City_Category'],drop_first=True)
# as the previous code is resulting in the true and false added the dtype as
↳INT.
df_city=pd.get_dummies(df['City_Category'],drop_first=True, dtype=int)
```

```
[ ]: df_city.head()
```

```
[ ]:  B  C
0  0  0
1  0  0
2  0  0
3  0  0
4  0  1
```

```
[ ]: df=pd.concat([df,df_city],axis=1)
df = df.drop('City_Category', axis =1)
```

```
df.head()
```

```
[ ]:  Product_ID  Gender  Age  Occupation  Stay_In_Current_City_Years  \
0  P00069042      0    1      10              2
1  P00248942      0    1      10              2
2  P00087842      0    1      10              2
3  P00085442      0    1      10              2
4  P00285442      1    7      16             4+

      Marital_Status  Product_Category_1  Product_Category_2  Product_Category_3  \
0                0                3                NaN                NaN
1                0                1                6.0             14.0
2                0               12                NaN                NaN
3                0               12               14.0                NaN
4                0                8                NaN                NaN

      Purchase  B  C
0      8370.0  0  0
1     15200.0  0  0
2      1422.0  0  0
3      1057.0  0  0
4      7969.0  0  1
```

```
[ ]: df.head()
```

```
[ ]:  Product_ID  Gender  Age  Occupation  Stay_In_Current_City_Years  \
0  P00069042      0    1      10              2
1  P00248942      0    1      10              2
2  P00087842      0    1      10              2
3  P00085442      0    1      10              2
4  P00285442      1    7      16             4+

      Marital_Status  Product_Category_1  Product_Category_2  Product_Category_3  \
0                0                3                NaN                NaN
1                0                1                6.0             14.0
2                0               12                NaN                NaN
3                0               12               14.0                NaN
4                0                8                NaN                NaN

      Purchase  B  C
0      8370.0  0  0
1     15200.0  0  0
2      1422.0  0  0
3      1057.0  0  0
4      7969.0  0  1
```

```
[ ]: ## Missing Values  
df.isnull().sum()
```

```
[ ]: Product_ID          0  
     Gender             0  
     Age                0  
     Occupation         0  
     Stay_In_Current_City_Years  0  
     Marital_Status     0  
     Product_Category_1  0  
     Product_Category_2 245982  
     Product_Category_3 545809  
     Purchase           233599  
     B                  0  
     C                  0  
     dtype: int64
```

```
[ ]: ## Focus on replacing missing values  
df['Product_Category_2'].unique()
```

```
[ ]: array([nan,  6., 14.,  2.,  8., 15., 16., 11.,  5.,  3.,  4., 12.,  9.,  
          10., 17., 13.,  7., 18.])
```

```
[ ]: df['Product_Category_2'].value_counts()
```

```
[ ]: Product_Category_2  
     8.0    91317  
     14.0   78834  
     2.0    70498  
     16.0   61687  
     15.0   54114  
     5.0    37165  
     4.0    36705  
     6.0    23575  
     11.0   20230  
     17.0   19104  
     13.0   15054  
     9.0     8177  
     12.0    7801  
     10.0    4420  
     3.0     4123  
     18.0    4027  
     7.0     854  
     Name: count, dtype: int64
```

```
[ ]: # For the categorical variables and discrete variables  
     # the Mode is used to replace the missing values.
```



```
df['Product_Category_2'].mode()[0]
#here the indexing is used to get the mode alone in the result.
```

```
[ ]: 8.0
```

```
[ ]: ## Replace the missing values with mode using the fillna
df['Product_Category_2']=df['Product_Category_2'].
    ↪fillna(df['Product_Category_2'].mode()[0])
```

```
[ ]: df['Product_Category_2'].isnull().sum()
```

```
[ ]: 0
```

```
[ ]: ## Product_category 3 replace missing values
df['Product_Category_3'].unique()
```

```
[ ]: array([nan, 14., 17.,  5.,  4., 16., 15.,  8.,  9., 13.,  6., 12.,  3.,
        18., 11., 10.] )
```

```
[ ]: df['Product_Category_3'].value_counts()
```

```
[ ]: Product_Category_3
16.0    46469
15.0    39968
14.0    26283
17.0    23818
 5.0     23799
 8.0     17861
 9.0     16532
12.0     13115
13.0      7849
 6.0      6888
18.0      6621
 4.0      2691
11.0      2585
10.0      2501
 3.0       878
Name: count, dtype: int64
```

```
[ ]: ## Replace the missing values with mode
df['Product_Category_3']=df['Product_Category_3'].
    ↪fillna(df['Product_Category_3'].mode()[0])
```

```
[ ]: df.head()
```

```
[ ]:   Product_ID  Gender  Age  Occupation  Stay_In_Current_City_Years  \
0  P00069042      0     1         10                             2
```

1	P00248942	0	1	10	2
2	P00087842	0	1	10	2
3	P00085442	0	1	10	2
4	P00285442	1	7	16	4+

	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	\
0	0	3	8.0	16.0	
1	0	1	6.0	14.0	
2	0	12	8.0	16.0	
3	0	12	14.0	16.0	
4	0	8	8.0	16.0	

	Purchase	B	C
0	8370.0	0	0
1	15200.0	0	0
2	1422.0	0	0
3	1057.0	0	0
4	7969.0	0	1

```
[ ]: df['Stay_In_Current_City_Years'].unique()
```

```
[ ]: array(['2', '4+', '3', '1', '0'], dtype=object)
```

```
[ ]: df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].str.  
      ↪replace('+','')
```

```
[ ]: df.head()
```

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	\
0	P00069042	0	1	10	2	
1	P00248942	0	1	10	2	
2	P00087842	0	1	10	2	
3	P00085442	0	1	10	2	
4	P00285442	1	7	16	4	

	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	\
0	0	3	8.0	16.0	
1	0	1	6.0	14.0	
2	0	12	8.0	16.0	
3	0	12	14.0	16.0	
4	0	8	8.0	16.0	

	Purchase	B	C
0	8370.0	0	0
1	15200.0	0	0
2	1422.0	0	0
3	1057.0	0	0

```
4    7969.0    0    1
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 783667 entries, 0 to 783666
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Product_ID                           783667 non-null  object
1   Gender                               783667 non-null  int64
2   Age                                  783667 non-null  int64
3   Occupation                           783667 non-null  int64
4   Stay_In_Current_City_Years          783667 non-null  object
5   Marital_Status                       783667 non-null  int64
6   Product_Category_1                  783667 non-null  int64
7   Product_Category_2                  783667 non-null  float64
8   Product_Category_3                  783667 non-null  float64
9   Purchase                            550068 non-null  float64
10  B                                    783667 non-null  int64
11  C                                    783667 non-null  int64
dtypes: float64(3), int64(7), object(2)
memory usage: 71.7+ MB
```


```
[ ]: ##convert object into integers
df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].
    .astype('int64')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 783667 entries, 0 to 783666
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Product_ID                           783667 non-null  object
1   Gender                               783667 non-null  int64
2   Age                                  783667 non-null  int64
3   Occupation                           783667 non-null  int64
4   Stay_In_Current_City_Years          783667 non-null  int64
5   Marital_Status                       783667 non-null  int64
6   Product_Category_1                  783667 non-null  int64
7   Product_Category_2                  783667 non-null  float64
8   Product_Category_3                  783667 non-null  float64
9   Purchase                            550068 non-null  float64
10  B                                    783667 non-null  int64
11  C                                    783667 non-null  int64
dtypes: float64(3), int64(8), object(1)
```

memory usage: 71.7+ MB

```
[ ]: ##Visualisation Age vs Purchased  
sns.barplot(x = 'Age',y = 'Purchase',hue='Gender',data=df)
```


```
[ ]: <Axes: xlabel='Age', ylabel='Purchase'>
```



## 2.1 Purchasing of men is high then women

```
[ ]: ## Visualization of Purchase with occupation  
sns.barplot(x = 'Occupation',y = 'Purchase',hue='Gender',data=df)
```

```
[ ]: <Axes: xlabel='Occupation', ylabel='Purchase'>
```




```
[ ]: plt.figure(figsize= (20,15))

plt.subplot(3,3,1)
sns.barplot(x = 'Product_Category_1',y = 'Purchase',hue='Gender',data=df)

plt.subplot(3,3,2)
sns.barplot(x = 'Product_Category_2',y = 'Purchase',hue='Gender',data=df)

plt.subplot(3,3,3)
sns.barplot(x = 'Product_Category_3',y = 'Purchase',hue='Gender',data=df)
plt.savefig('product-wise purchasing.jpg')
```



```
[ ]: #sns.pairplot(df, hue = 'Product_ID', diag_kind = 'kde')
#plt.savefig('pairplot.jpg')
```

```
[ ]: ##Feature Scaling
# test data is taken where the purchase is null as we compbined before
# train is taken with purchase which are not Nan
df_test=df[df['Purchase'].isnull()]
```

```
[ ]: df_test.head(5)
```

```
[ ]:      Product_ID  Gender  Age  Occupation  Stay_In_Current_City_Years  \
550068  P00128942      1    5         7             2
550069  P00113442      1    3        17             0
550070  P00288442      0    4         1             4
550071  P00145342      0    4         1             4
550072  P00053842      0    3         1             1
```

```
      Marital_Status  Product_Category_1  Product_Category_2  \
550068              1                  1             11.0
550069              0                  3              5.0
550070              1                  5             14.0
550071              1                  4              9.0
550072              0                  4              5.0
```

```
      Product_Category_3  Purchase  B  C
550068              16.0      NaN  1  0
550069              16.0      NaN  0  1
550070              16.0      NaN  1  0
550071              16.0      NaN  1  0
550072              12.0      NaN  0  1
```

```
[ ]: df_train=df[~df['Purchase'].isnull()] # ~ not having
df_train.head(5)
```

```
[ ]:      Product_ID  Gender  Age  Occupation  Stay_In_Current_City_Years  \
0  P00069042      0    1         10             2
1  P00248942      0    1         10             2
2  P00087842      0    1         10             2
3  P00085442      0    1         10             2
4  P00285442      1    7         16             4
```

```
      Marital_Status  Product_Category_1  Product_Category_2  Product_Category_3  \
0              0                  3             8.0             16.0
1              0                  1             6.0             14.0
2              0                 12             8.0             16.0
3              0                 12            14.0             16.0
```

4	0	8	8.0	16.0
---	---	---	-----	------

	Purchase	B	C
0	8370.0	0	0
1	15200.0	0	0
2	1422.0	0	0
3	1057.0	0	0
4	7969.0	0	1

```
[ ]: X=df_train.drop(['Purchase', 'Product_ID'],axis=1)
X.head()
```

```
[ ]:      Gender  Age  Occupation  Stay_In_Current_City_Years  Marital_Status  \
0         0     1         10                2                0
1         0     1         10                2                0
2         0     1         10                2                0
3         0     1         10                2                0
4         1     7         16                4                0
```

	Product_Category_1	Product_Category_2	Product_Category_3	B	C
0	3	8.0	16.0	0	0
1	1	6.0	14.0	0	0
2	12	8.0	16.0	0	0
3	12	14.0	16.0	0	0
4	8	8.0	16.0	0	1

```
[ ]: X.shape
```

```
[ ]: (550068, 10)
```

```
[ ]: y=df_train['Purchase']
```

```
[ ]: y.shape
```

```
[ ]: (550068,)
```

```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=42)
```

```
[ ]: X_train.head()
```

```
[ ]:      Gender  Age  Occupation  Stay_In_Current_City_Years  Marital_Status  \
396876      1     2         14                3                0
433826      1     6          0                0                1
516298      1     4         17                0                0
193380      1     3          4                1                0
```

273542	0	4	20	3	1
	Product_Category_1	Product_Category_2	Product_Category_3	B	C
396876	1	2.0	16.0	1	0
433826	8	16.0	16.0	0	0
516298	3	4.0	12.0	0	1
193380	8	16.0	16.0	1	0
273542	3	4.0	12.0	1	0

```
[ ]: ## feature Scaling
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train =sc.fit_transform(X_train)
X_test =sc.transform(X_test)
```

```
[ ]: from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
```

```
[ ]: lr = LinearRegression()
lr.fit(X_train, y_train)
lr.score(X_train, y_train)*100, lr.score(X_test,y_test)*100
```

```
[ ]: (13.210555628187514, 12.948768560712399)
```

```
[ ]: dt = DecisionTreeRegressor()
dt.fit(X_train, y_train)
dt.score(X_train, y_train)*100, dt.score(X_test,y_test)*100
```

```
[ ]: (79.83631647792538, 55.82640534204286)
```

```
[ ]: #lr3 = Lasso()
#lr3.fit(X_train, y_train)
#lr3.score(X_train, y_train)*100, lr3.score(X_test,y_test)*100
```

```
[ ]: lr4 = Ridge()
lr4.fit(X_train, y_train)
lr4.score(X_train, y_train)*100, lr4.score(X_test,y_test)*100
```

```
[ ]: (13.210555628099762, 12.948769558874373)
```

```
[ ]: rfr = RandomForestRegressor()
rfr.fit(X_train, y_train)
rfr.score(X_train, y_train)*100, rfr.score(X_test,y_test)*100
```



```
[ ]: (78.75051863025938, 62.922348370441014)
```

```
[ ]: knr = KNeighborsRegressor()  
knr.fit(X_train, y_train)  
knr.score(X_train, y_train)*100, knr.score(X_test, y_test)*100
```

```
[ ]: (64.73951731332154, 50.659057532616615)
```

```
[ ]: from sklearn.metrics import r2_score
```

```
[ ]: print(r2_score(y_test, rfr.predict(X_test)))
```

```
0.6292234837044102
```

```
[ ]: from sklearn.metrics import classification_report
```

```
[ ]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[ ]: gbr = GradientBoostingRegressor()
```

```
[ ]: gbr.fit(X_train, y_train)  
print(r2_score(y_test, gbr.predict(X_test)))
```

```
0.6444959838731033
```

```
[ ]: gbr.score(X_train, y_train)*100, gbr.score(X_test, y_test)*100
```

```
[ ]: (64.96087814238759, 64.44959838731033)
```

```
[ ]: from xgboost import XGBRegressor, XGBRFRegressor  
xgb = XGBRegressor()  
xgbr = XGBRFRegressor()
```

```
[ ]: xgb.fit(X_train, y_train)  
print(r2_score(y_test, xgb.predict(X_test)))  
  
xgbr.fit(X_train, y_train)  
print(r2_score(y_test, xgbr.predict(X_test)))
```

```
0.6668955124144831
```

```
0.5845697344685554
```

```
[ ]: xgb.score(X_train, y_train)*100, xgb.score(X_test, y_test)*100
```

```
[ ]: (67.99197211448806, 66.68955124144831)
```

```
[ ]: xgbr.score(X_train, y_train)*100, xgbr.score(X_test, y_test)*100
```

```
[ ]: (58.9592269938375, 58.45697344685554)
```

```
[ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
[ ]: def evaluate_model(true, predicted):  
    mae = mean_absolute_error(true, predicted)  
    mse = mean_squared_error(true, predicted)  
    rmse = np.sqrt(mean_squared_error(true, predicted))  
    r2_square = r2_score(true, predicted)  
    return mae, mse, rmse, r2_square
```

```
[ ]: from sklearn.ensemble import AdaBoostRegressor
```

```
[ ]: models = {  
    "K-Neighbors Regressor": KNeighborsRegressor(),  
    "Random Forest Regressor": RandomForestRegressor(),  
    "XGBRegressor": XGBRegressor(),  
    "AdaBoost Regressor": AdaBoostRegressor()  
}
```

```
[ ]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV  
from sklearn.metrics import make_scorer  
  
# Define hyperparameter ranges for each model  
param_grid = {  
    "K-Neighbors Regressor": {"n_neighbors": [3, 5, 7]},  
    "Random Forest Regressor": {'n_estimators': [8,16,32,64,128,256]},  
    ↪ "max_depth": [3, 5, 7]},  
    "XGBRegressor": {'depth': [6,8,10], 'learning_rate': [0.01, 0.05, 0.  
    ↪ 1], 'iterations': [30, 50, 100]},  
    "AdaBoost Regressor": {'learning_rate': [.1, .01, 0.5, .001], 'n_estimators':  
    ↪ [8,16,32,64,128,256]}  
}  
  
model_list = []  
r2_list = []  
  
for model_name, model in models.items():  
    # Create a scorer object to use in grid search  
    scorer = make_scorer(r2_score)  
  
    # Perform grid search to find the best hyperparameters  
    grid_search = GridSearchCV(  
        model,  
        param_grid[model_name],  
        scoring=scorer,  
        cv=5,
```

```

        n_jobs=-1
    )

    grid_search.fit(X_train, y_train) # Make predictions

    y_train_pred = grid_search.predict(X_train)

    y_test_pred = grid_search.predict(X_test)

    # Evaluate Train and Test dataset
    model_train_mae, model_train_mse, model_train_rmse, model_train_r2 = _
    evaluate_model(y_train, y_train_pred)

    model_test_mae, model_test_mse, model_test_rmse, model_test_r2 = _
    evaluate_model(y_test, y_test_pred)

    print(model_name)

    model_list.append(model_name)

    print('Best hyperparameters:', grid_search.best_params_)
    print('Model performance for Training set')
    print("- Root Mean Squared Error: {:.4f}".format(model_train_rmse))
    print("- Mean Squared Error: {:.4f}".format(model_train_mse))
    print("- Mean Absolute Error: {:.4f}".format(model_train_mae))
    print("- R2 Score: {:.4f}".format(model_train_r2))
    print('-----')
    print('Model performance for Test set')
    print("- Root Mean Squared Error: {:.4f}".format(model_test_rmse))
    print("- Mean Squared Error: {:.4f}".format(model_test_mse))
    print("- Mean Absolute Error: {:.4f}".format(model_test_mae))
    print("- R2 Score: {:.4f}".format(model_test_r2))
    r2_list.append(model_test_r2)
    print('='*35)
    print('\n')

```

K-Neighbors Regressor

Best hyperparameters: {'n\_neighbors': 5}

Model performance for Training set

- Root Mean Squared Error: 2983.5495
- Mean Squared Error: 8901567.7598
- Mean Absolute Error: 2192.1820
- R2 Score: 0.6474

-----

Model performance for Test set

- Root Mean Squared Error: 3526.3571
- Mean Squared Error: 3526.3571

```
- Mean Absolute Error: 2564.8292
- R2 Score: 0.5066
=====
```

Random Forest Regressor

Best hyperparameters: {'max\_depth': 7, 'n\_estimators': 128}

Model performance for Training set

```
- Root Mean Squared Error: 2964.1804
- Mean Squared Error: 8786365.4030
- Mean Absolute Error: 2236.1604
- R2 Score: 0.6520
-----
```

Model performance for Test set

```
- Root Mean Squared Error: 2987.6687
- Mean Squared Error: 2987.6687
- Mean Absolute Error: 2251.5360
- R2 Score: 0.6458
=====
```

```
/usr/local/lib/python3.10/dist-
packages/joblib/externals/loky/process_executor.py:752: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning:
[11:26:33] WARNING: /workspace/src/learner.cc:740:
Parameters: { "depth", "iterations" } are not used.
```

```
warnings.warn(msg, UserWarning)
```

XGBRegressor

Best hyperparameters: {'depth': 6, 'iterations': 30, 'learning\_rate': 0.1}

Model performance for Training set

```
- Root Mean Squared Error: 2895.8257
- Mean Squared Error: 8385806.4002
- Mean Absolute Error: 2178.0139
- R2 Score: 0.6678
-----
```

Model performance for Test set

```
- Root Mean Squared Error: 2928.4398
- Mean Squared Error: 2928.4398
- Mean Absolute Error: 2199.5659
- R2 Score: 0.6597
=====
```

[ ]: